

CS222: Systems Programming

Memory Management

February 19th, 2008

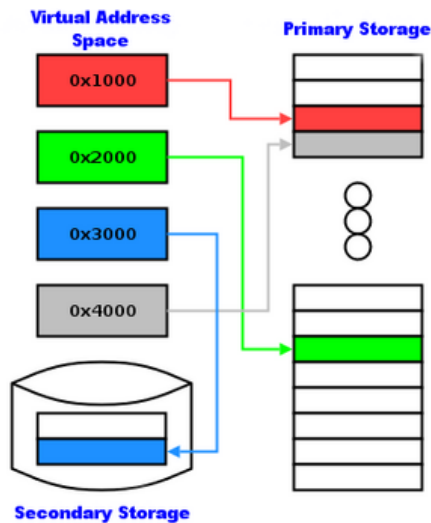
Last Class

- Error Handling
 - Exception Handling
 - Console Control Handlers
 - Vectored Exception Handling

Today's Class

- Memory management
 - Overview
 - Heap management
 - Memory-mapped files
 - Dynamic link libraries

Memory Management I



Process and Memory Space

- Each process has **its own virtual address space**
 - Up to 4 GB of memory (32-bit)
 - Actually 2GB (3GB possible)
- All **threads of a process** can access its virtual address space
 - However, they **cannot access memory** that belongs to **another process**

Virtual Address Space

- Virtual address of a process **does not represent the actual physical location** of an object in memory
- Each process maintains its **page map**
 - Internal data structure used to **translate virtual addresses into corresponding physical addresses**
 - Each time a thread references an address, the system translates the virtual address to physical address

Virtual and Physical Memory

- Virtual address space of a process **can be smaller or larger** than the total physical memory available on the computer
- The subset of the virtual address space of a process that resides in physical memory is called **working set**
 - If the threads of a process **attempt to use more** physical memory than is currently available, then **the system pages** some memory contents to disk

Pages

- A page is **a unit of memory**, into which physical storage and the virtual address space of each process are organized
 - Size depends on the host computer
- When a page is moved in physical memory, the system updates the page maps of the affected processes
- When the system needs space in physical memory, it moves the least recently used pages of physical memory to the paging file

Page State

- The pages of a process's virtual address space can be in one of the following states
 - Free
 - Neither committed nor reserved, but available
 - Not accessible to the process
 - Attempting to read from or write to a free page results in **access violation exception**
 - `VirtualFree` or `VirtualFreeEx`

Page State, cont

– Reserved

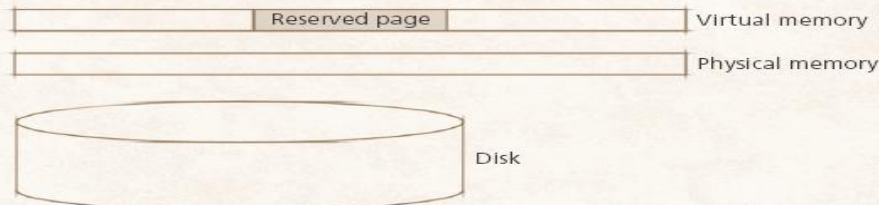
- Reserved for future use
- Address range cannot be used by other allocation functions
- Not accessible and has no physical storage associated with it
- Available to be committed
- `VirtualAlloc` or `VirtualAllocEx`

– Committed

- Physical storage is allocated, and access is controlled
- When process terminates, it is released
- `VirtualAlloc` or `VirtualAllocEx`

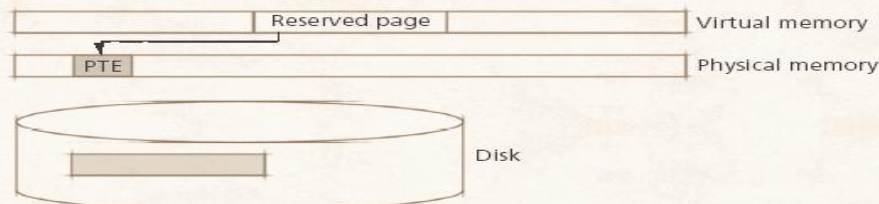
Page State, cont

a) Reserve



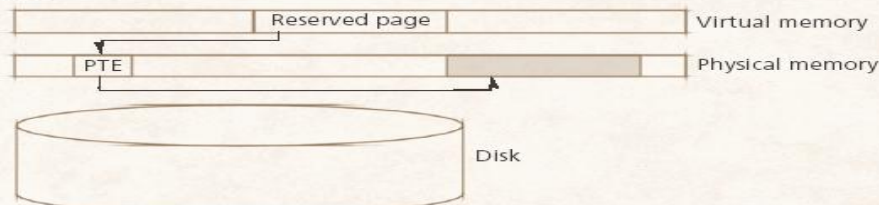
First, a process reserves memory. The VMM allocates space for the requested memory in the process's virtual address space.

b) Commit



Next, the process commits the reserved memory. The VMM allocates a page table entry (PTE) and ensures that it can allocate space in a pagefile on disk.

c) Access



Finally, the process accesses the committed memory. The VMM writes the data to a zeroed page in main memory and sets the page table entry (PTE) to point to this page.

Scope of Allocated Memory

- All memory allocated by **memory allocation functions**
 - Is **process-wide**
 - `HeapAlloc`, `VirtualAlloc`, `GlobalAlloc`, `LocalAlloc`
- All memory allocated by a DLL is allocated in the **address space of the process that called the DLL**
- In order to create shared memory, we must use file mapping

Page Faults

- References to pages **not in memory**
 - Most virtual pages will not be in physical memory
 - OS loads the data from disk, either from
 - System swap file, or
 - Normal file
- For performance purpose, programs should be designed minimize page faults

GetSystemInfo

- A function returning information about the current system
 - `SYSTEM_INFO` structure contains information including the architecture and type of a processor, the number of processors, page size, etc

```
VOID GetSystemInfo(  
    LPSYSTEM_INFO lpSystemInf);
```

Example: GetSystemInfo

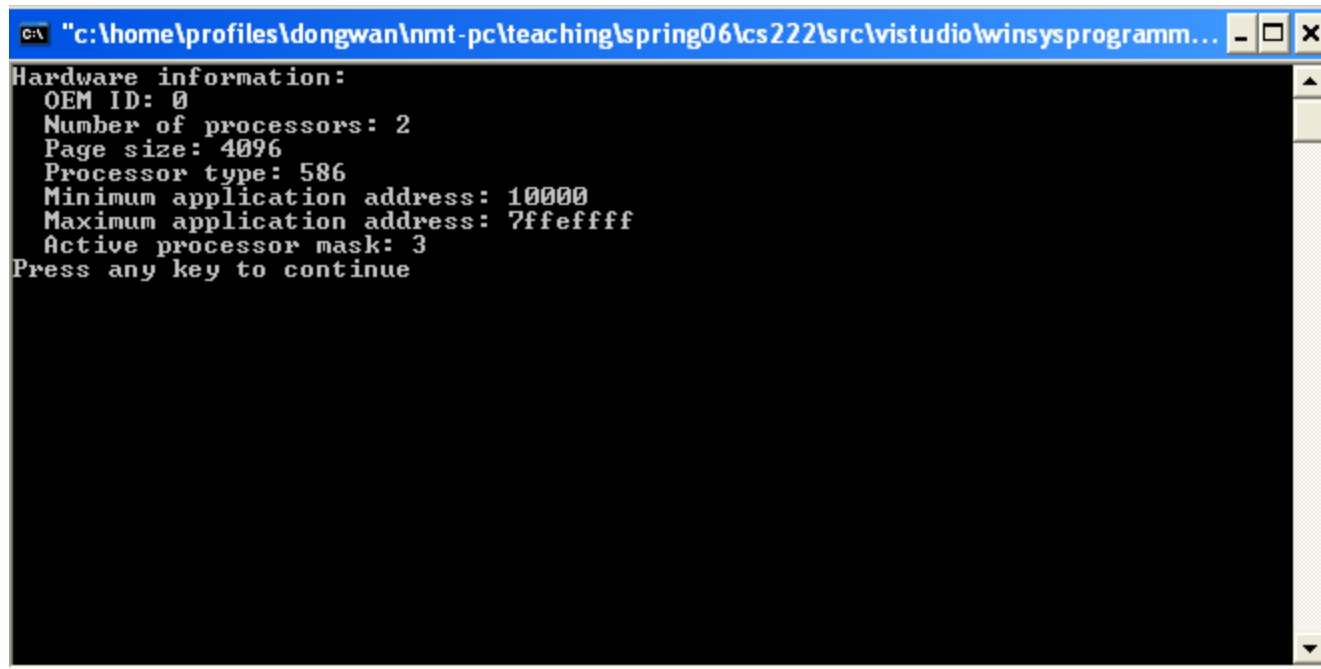
```
void main()
{
    SYSTEM_INFO siSysInfo;

    GetSystemInfo(&siSysInfo);

    // Display the contents of the SYSTEM_INFO structure.

    printf("Hardware information: \n");
    printf("  OEM ID: %u\n", siSysInfo.dwOemId);
    printf("  Number of processors: %u\n",
        siSysInfo.dwNumberOfProcessors);
    printf("  Page size: %u\n", siSysInfo.dwPageSize);
    printf("  Processor type: %u\n", siSysInfo.dwProcessorType);
    printf("  Minimum application address: %lx\n",
        siSysInfo.lpMinimumApplicationAddress);
    printf("  Maximum application address: %lx\n",
        siSysInfo.lpMaximumApplicationAddress);
    printf("  Active processor mask: %u\n",
        siSysInfo.dwActiveProcessorMask);
}
```

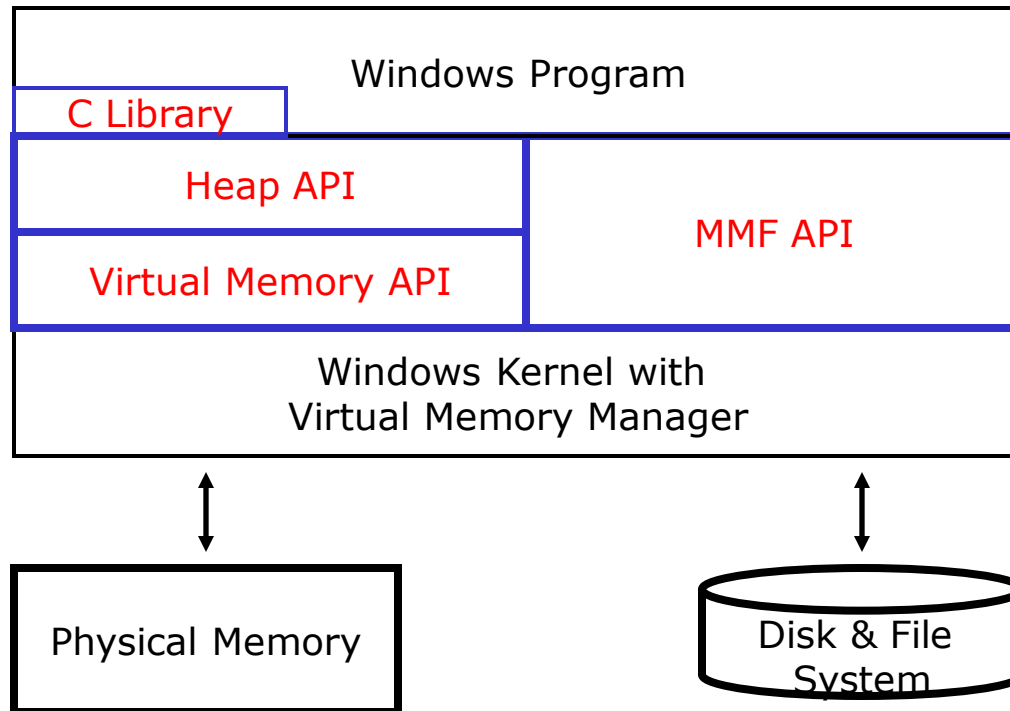
Example: GetSystemInfo



A screenshot of a Windows command prompt window. The title bar shows the path "c:\home\profiles\dongwan\mnt-pc\teaching\spring06\cs222\src\vstudio\winsysprogramm...". The command prompt displays the following text:

```
Hardware information:  
OEM ID: 0  
Number of processors: 2  
Page size: 4096  
Processor type: 586  
Minimum application address: 10000  
Maximum application address: 7ffeffff  
Active processor mask: 3  
Press any key to continue
```


Windows Memory Management



Heaps

- A heap is used for **allocating and freeing objects dynamically** for use by the program. Heap operations are called for when
 - The number and size of objects needed by the program are not known ahead of time
 - An object is too large to fit into a stack allocator

Heap Management

- A process **can contain several heaps** for following reasons
 - Fairness
 - Multithreaded performance
 - Allocation efficiency
 - Deallocation efficiency
 - Locality of reference efficiency
- Often a single heap is sufficient. In that case, use the C library memory management functions
 - `malloc`, `calloc`, `realloc`, `free`, etc

GetProcessHeap

- A function used for obtaining a handle to the heap of the calling process
 - Heap handle is **necessary** when you are allocating memory
 - Each process has its own **default heap, which is used by malloc**

```
HANDLE GetProcessHeap( VOID );
```

Return: The handle for the process's heap: NULL on failure

HeapCreate

- A function used for creating a heap object that can be used by the calling process
 - Reserve space in the virtual address space of the process
 - Allocate physical storage for a specified initial portion
 - `flOptions`
 - `HEAP_GENERATE_EXCEPTIONS`
 - `HEAP_NO_SERIALIZE`

```
HANDLE HeapCreate(  
    DWORD flOptions,  
    SIZE_T dwInitialSize,  
    SIZE_T dwMaximumSize);
```

Return: The handle for the heap: NULL on failure

HeapDestroy

- A function used for destroying an entire heap
 - Decommit and release all the pages of a private heap object
 - Be careful **not to destroy the process's heap**
- Destroying a heap is a quick way to free data structures without traversing them to delete one element at a time

```
BOOL HeapDestroy( HANDLE hHeap );
```

HeapAlloc

- A function used for allocating a block of memory from a heap
 - dwFlags
 - HEAP_GENERATE_EXCEPTIONS
 - HEAP_NO_SERIALIZE
 - HEAP_ZERO_MEMORY
- Use HeapFree function to deallocate memory

```
LPVOID HeapAlloc(  
    HANDLE hHeap,  
    DWORD dwFlags,  
    SIZE_T dwBytes);
```

Return: A pointer to the allocated memory block, or NULL on failure

HeapReAlloc

- A function used for reallocating a block of memory from a heap

```
LPVOID HeapReAlloc(  
    HANDLE hHeap,  
    DWORD dwFlags,  
    LPVOID lpMem  
    SIZE_T dwBytes);
```

Return: A pointer to the reallocated memory block, or NULL on failure

HEAP_NO_SERIALIZE

- Use for small performance gain
- Requirements
 - No multi-threaded programming
or
 - Each thread uses its own heap
or
 - Program has its own mutual exclusion mechanism

Summary: Heap Management

- The normal process for using heaps is as follows
 1. Get a **heap handle** with either `HeapCreate` or `GetProcessHeap`
 2. **Allocate blocks** within the heap using `HeapAlloc`
 3. Optionally, free some or all of the individual blocks with `HeapFree`
 4. **Destroy** the heap and close the handle with `HeapDestroy`

Example: HeapCreate/HeapAlloc

```
HANDLE hHeap;
SIZE_T nBufferSize;

...

/* allocate memory for the buffer */
__try{
    hHeap = HeapCreate(HEAP_GENERATE_EXCEPTIONS | HEAP_NO_SERIALIZE,
                      nBufferSize, 0);    // growable heap size

    cBuffer = HeapAlloc(hHeap, HEAP_ZERO_MEMORY, sizeof(TCHAR)*nBufferSize);
}

__except(EXCEPTION_EXECUTE_HANDLER){
    printf("Exception occurred... : %x", GetExceptionCode());
}

...

/* free allocated memory */
HeapDestroy(hHeap);
```

Review

- Memory management
 - Overview
 - Heap management
 - Memory-mapped files
 - Dynamic link libraries
- Recommended reading for next class
 - Chapter 6 in Windows System Programming

Next Class

- Quiz
- Homework due **next Tuesday**