

CS222: Systems Programming

Memory Management II

February 21st, 2008

Last Class

- Memory Management
 - Overview
 - Heap management
 - Memory-mapped files
 - Dynamic link libraries

Today's Class

- **Memory Management**
 - Overview
 - Heap management
 - **Memory-mapped files**
 - Dynamic link libraries

Heap Review

- Memory reserved by HeapCreate is *not necessarily contiguous*
- Memory allocated by HeapAlloc is *contiguous*

GetProcessHeap

- A function used for obtaining a handle to the heap of the calling process
 - Heap handle is **necessary** when you are allocating memory
 - Each process has its own **default heap, which is used by malloc**

```
HANDLE GetProcessHeap( VOID );
```

Return: The handle for the process's heap: NULL on failure

HeapCreate

- A function used for creating a heap object that can be used by the calling process
 - Reserve space in the virtual address space of the process
 - Allocate physical storage for a specified initial portion
 - `flOptions`
 - `HEAP_GENERATE_EXCEPTIONS`
 - `HEAP_NO_SERIALIZE`

```
HANDLE HeapCreate(  
    DWORD flOptions,  
    SIZE_T dwInitialSize,  
    SIZE_T dwMaximumSize);
```

Return: The handle for the heap: NULL on failure

HeapDestroy

- A function used for destroying an entire heap
 - Decommit and release all the pages of a private heap object
 - Be careful **not to destroy the process's heap**
- Destroying a heap is a quick way to free data structures without traversing them to delete one element at a time

```
BOOL HeapDestroy( HANDLE hHeap );
```

HeapAlloc

- A function used for allocating a block of memory from a heap
 - dwFlags
 - HEAP_GENERATE_EXCEPTIONS
 - HEAP_NO_SERIALIZE
 - HEAP_ZERO_MEMORY
- Use HeapFree function to deallocate memory

```
LPVOID HeapAlloc(  
    HANDLE hHeap,  
    DWORD dwFlags,  
    SIZE_T dwBytes);
```

Return: A pointer to the allocated memory block, or NULL on failure

HeapReAlloc

- A function used for reallocating a block of memory from a heap

```
LPVOID HeapReAlloc(  
    HANDLE hHeap,  
    DWORD dwFlags,  
    LPVOID lpMem  
    SIZE_T dwBytes);
```

Return: A pointer to the reallocated memory block, or NULL on failure

HEAP_NO_SERIALIZE

- Use for small performance gain
- Requirements
 - No multi-threaded programming
or
 - Each thread uses its own heap
or
 - Program has its own mutual exclusion mechanism

Summary: Heap Management

- The normal process for using heaps is as follows
 1. Get a **heap handle** with either `HeapCreate` or `GetProcessHeap`
 2. **Allocate blocks** within the heap using `HeapAlloc`
 3. Optionally, free some or all of the individual blocks with `HeapFree`
 4. **Destroy** the heap and close the handle with `HeapDestroy`

Memory-mapped Files

- Memory-mapped file functionality
 - Map **virtual memory space** directly **to normal files**
- Advantages
 - No need to perform direct file I/O
 - Data structures created in memory will be saved in the file for later use
 - In-memory algorithms can process file data even though the file is much larger than available physical memory
 - Improvement of file processing performance
 - No need to manage buffers and the file data
 - **Multiple processes can share memory**

File Mapping Objects

- As **the first step** to use MMF, we
 - Need to create a file mapping object on an open file
 - Can give names to the object so that it can be accessible to other processes for shared memory
 - Can protect the object using security attributes
- Use `CreateFileMapping` function

CreateFileMapping

- A function used for creating or opening a named or unnamed file mapping object for specified file

```
HANDLE CreateFileMapping (  
    HANDLE hFile,  
    LPSECURITY_ATTRIBUTE lpsa,  
    DWORD dwProtect,  
    DWORD dwMaximumSizeHigh,  
    DWORD dwMaximumSizeLow,  
    LPCTSTR lpMapName  
);
```

Return: If function succeeds, the return value is a handle.
Otherwise, the return value is NULL

Example: CreateFileMapping

```
hMap = CreateFileMapping(...);  
  
if (hMap != NULL && GetLastError() == ERROR_ALREADY_EXISTS)  
{  
    CloseHandle(hMap);  
    hMap = NULL;  
}  
  
return hMap;
```

OpenFileMapping

- A function used for opening a named file mapping object

```
HANDLE OpenFileMapping (  
    DWORD dwDesiredAccess,  
    BOOL hInheritHandle,  
    LPCTSTR lpName  
);
```

Return: If function succeeds, the return value is a handle.
Otherwise, the return value is NULL

Mapping Address to Object

- As **the second step**, we
 - Need to allocate virtual memory space and map it to a file through the mapping object
 - Similar to `HeapAlloc`
 - Much coarser – larger allocation units
- Use `MapViewOfFile` function

MapViewOfFile

- A function used for mapping a view of a file mapping into the address space of a calling process

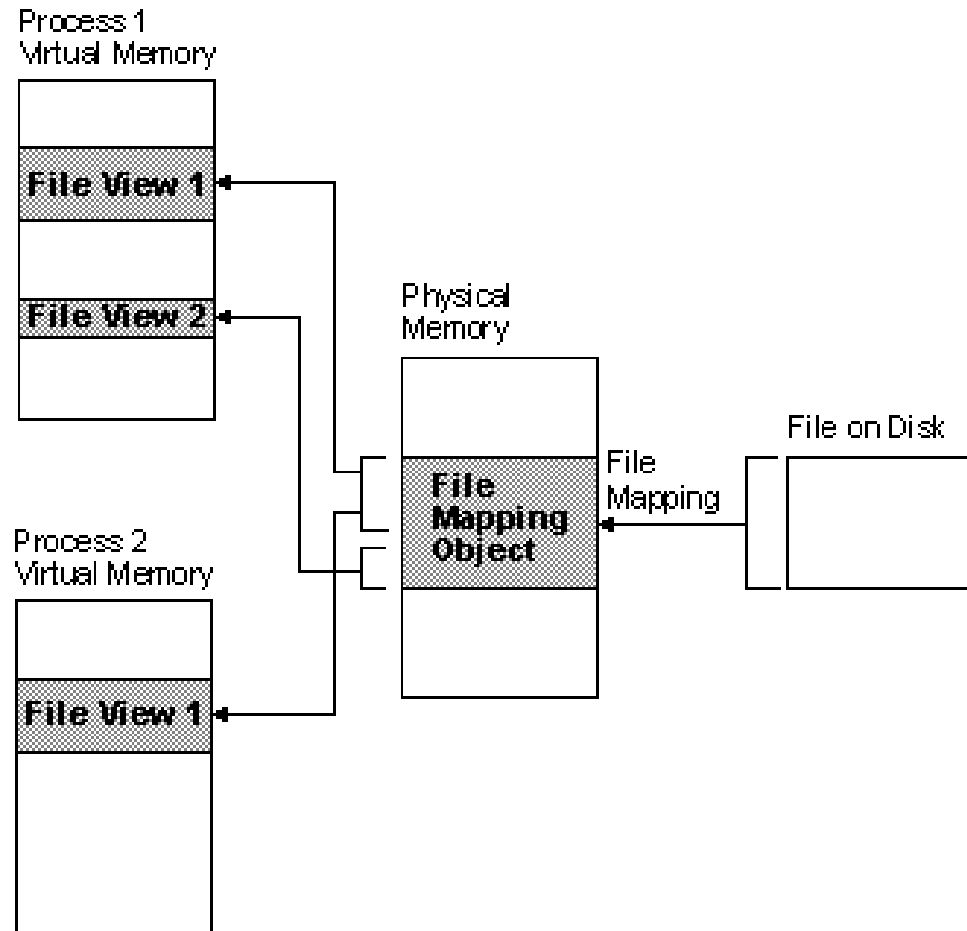
```
LPVOID MapViewOfFile (  
    HANDLE hFileMappingObject,  
    DWORD dwDesiredAccess,  
    DWORD dwFileOffsetHigh,  
    DWORD dwFileOffsetLow,  
    SIZE_T dwNumberOfBytesToMap  
);
```

Return: If function succeeds, the return value is starting address of the mapped view. Otherwise, NULL

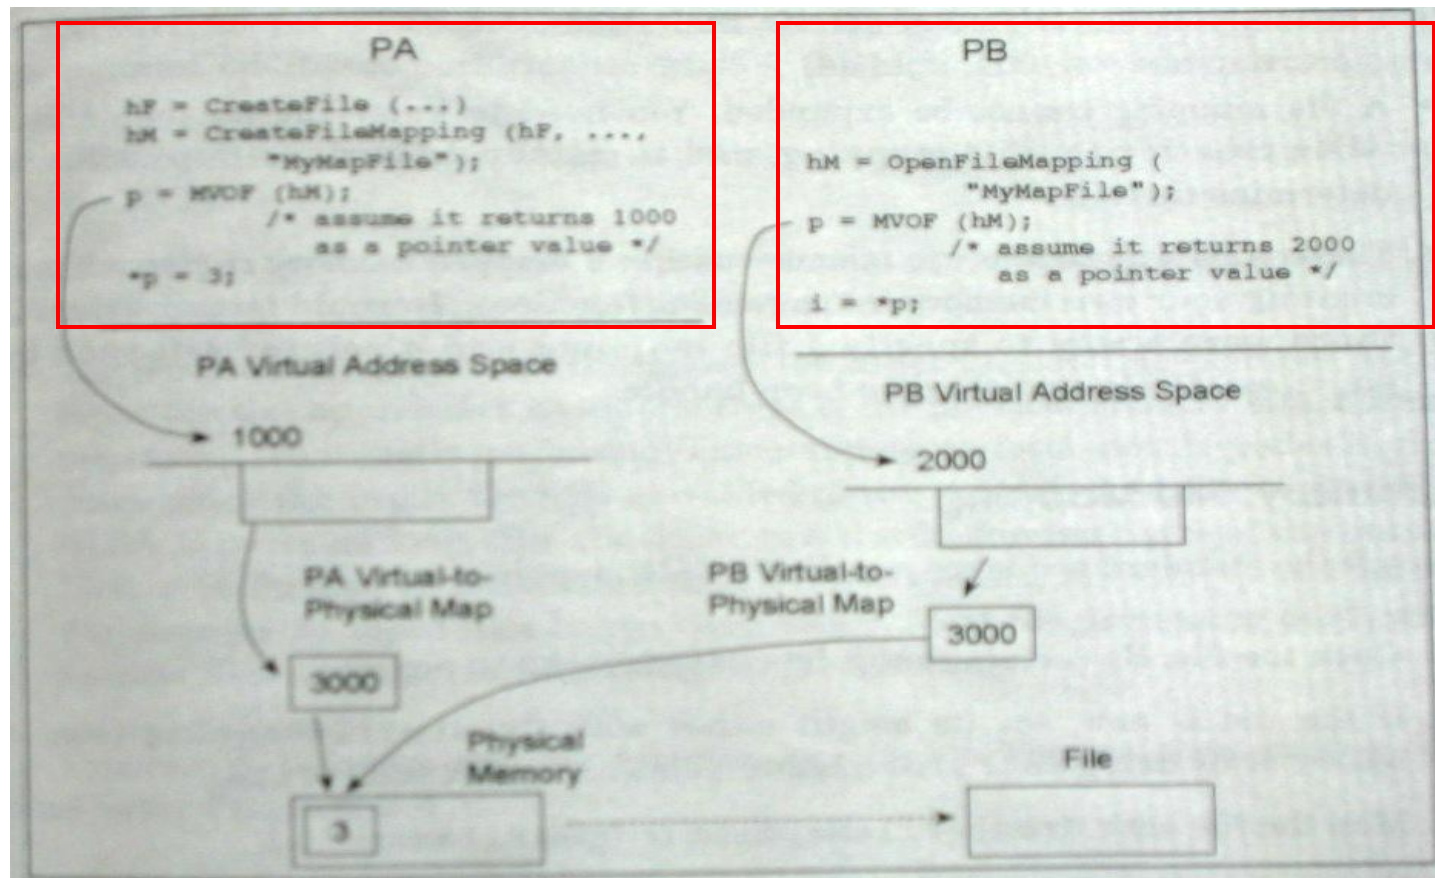
MapViewOfFile

- `MapViewOfFileEx` is similar
 - Must specify a starting memory address
- Use `UnmapViewOfFile` to release memory

Addresses Mapped to a File



Shared Memory



FlushViewOfFile

- Forces system to write changed pages to disk
 - Can be troublesome when concurrently using regular I/O
 - ReadFile
 - WriteFile
 - Coherency is not ensured
 - Do not use these I/O functions for mapped files

Example: MMF-P1

```
#include <windows.h>
#include <stdio.h>

#define BUF_SIZE 256
TCHAR szName[]=TEXT("MyFileMappingObject");
TCHAR szMsg[]=TEXT("Message from first process");

void main(){
    HANDLE hMapFile;
    LPCTSTR pBuf;

    hMapFile = CreateFileMapping(
        INVALID_HANDLE_VALUE,    // use paging file
        NULL,                    // default security
        PAGE_READWRITE,          // read/write access
        0,                        // max. object size
        BUF_SIZE,                 // buffer size
        szName);                  // name of mapping object
    if (hMapFile == NULL || hMapFile == INVALID_HANDLE_VALUE) {
        printf("Could not create file mapping object (%d).\n",
            GetLastError());
        return;
    }
}
```

Example: MMF-P1

```
pBuf = (LPTSTR) MapViewOfFile(hMapFile,    // handle to map object
                             FILE_MAP_ALL_ACCESS, // read/write permission
                             0,
                             0,
                             BUF_SIZE);

if (pBuf == NULL)
{
    printf("Could not map view of file (%d).\n",
          GetLastError());
    return;
}

CopyMemory((PVOID)pBuf, szMsg, strlen(szMsg));
getch();

UnmapViewOfFile(pBuf);

CloseHandle(hMapFile);
}
```


Example: MMF-P2

```
#include <windows.h>
#include <stdio.h>
#include <conio.h>

#define BUF_SIZE 256
TCHAR szName[]=TEXT("MyFileMappingObject");

void main()
{
    HANDLE hMapFile;
    LPCTSTR pBuf;

    hMapFile = OpenFileMapping(
        FILE_MAP_ALL_ACCESS,    // read/write access
        FALSE,                  // do not inherit the name
        szName);                // name of mapping object

    if (hMapFile == NULL)
    {
        printf("Could not open file mapping object (%d).\n",
            GetLastError());
        return;
    }
}
```

Example: MMF-P2

```
pBuf = MapViewOfFile(hMapFile,    // handle to mapping object
                    FILE_MAP_ALL_ACCESS, // read/write permission
                    0,
                    0,
                    BUF_SIZE);


if (pBuf == NULL)
{
    printf("Could not map view of file (%d).\n",
          GetLastError());
    return;
}

MessageBox(NULL, pBuf, TEXT("Process2"), MB_OK);

UnmapViewOfFile(pBuf);

CloseHandle(hMapFile);
}
```

Result



File Mapping Limitation

- File mapping is powerful and useful, however
 - File mapping **cannot be expanded**
 - **No way to allocate memory** within a mapped memory region

Summary: MMF

- Standard sequence required to use MMF
 1. Open the file
 2. If the file is new, set the file length either with `CreateFileMapping` or by using `SetFilePointer` followed by `SetEndOfFile`
 3. Map the file with `CreateFileMapping` or `OpenFileMapping`
 4. Create one or more views with `MapViewOfFile`
 5. Access the file through memory references
 6. On completion, un-map the file and close handles for file mapping object as well as file

Review

- Memory management
 - Overview
 - Heap management
 - Memory-mapped files
 - Dynamic link libraries
- Recommended reading for next class
 - Chapter 6 in Windows System Programming